

INTERACTIONS AND OTHER EXPANDED FEATURES

Contents

1	Learning objectives	1
2	Dealing with categorical data	1
3	Interaction	1
3.1	Interactions	3
4	Non-linear models	4
5	Orthogonality and Fourier analysis	6

1 Learning objectives

- Understand how to add interactions to linear regression models
- Be able to interpret regression coefficients when interactions are present
- Understand how to make residual plots and especially why we plot the predicted value of the response variable on the x axis.
- Understand how to build nonlinear models by adding “features”, and why for continuous predictors sines and cosines make good features.

2 Dealing with categorical data

- One situation in which models with multiple predictors frequently arises is when trying to predict a Y variable based on categorical predictors, such as race. In this case, we need to transform the categories into numerical values. For example, if there are two categories (e.g. YES and NO) we map our variable to 0 or 1. If we have 3 categories (e.g. White, Black, Other), we might first think to map them to 0, 1 and 2. This has a problem though: A change from 1 to 2 should not necessarily correspond to a change from 0 to 1. In other words, **there is no clear ordering of the x values**. Sometimes we refer to such predictors and qualitative rather than quantitative, since they express a quality of our data points instead of a numerical quantity.
- To address this issue, we create dummy variables. In particular, In order to take a categorical variable and transform it into a set of indicator variables in python, we use the python function `get_dummies`. The usage of this is illustrated in the following example.

3 Interaction

- The important assumption of the multiple predictor regression models we have seen so far is that the “effect” of one predictor does not depend on the value of the other. Here are some examples where this could be violated:
 - The difference in test scores between kids whose mothers did and did not go to high school depends on their mother’s score on the iq test.
 - The association between earnings and height depends on gender, e.g. being taller tends to give men a larger advantage than women
 - The effect of a drug for treating covid depends on whether someone has had covid before.

We call the dependencies described above interactions. It turns out it is possible to include interactions within the regression modeling framework we have already introduced, as is illustrated by the following example.

- Let's work in the case of two predictors. If you'd like, you can refer to Example 2 to make this math more concrete and always replace X_1 with X_{hs} and X_2 with X_{iq} . In each of the cases above, what is going on is that the slope of $\mathbb{E}[Y|X_1, X_2]$ vs. X_1 for fixed X_2 is not a constant, β_1 , but rather something that depends on X_2 . The simplest way to account for this is to let effect of X_2 on the slope be linear in X_2 , so instead of β_1 being the slope of $\mathbb{E}[Y|X_1, X_2]$ vs. X_1 we would have $\beta_1 + \beta_{1,2}X_2$ be this slope, but this leads to a regression model with a nonlinear term:

$$Y = \beta_0 + (\beta_1 + \beta_{1,2}X_2)X_1 + \beta_2X_2 = \beta_0 + \beta_1X_1 + \beta_2X_2 + \beta_{1,2}X_1X_2 + \epsilon$$

We call X_1X_2 an interaction term, but Yuck! This isn't linear anymore and the regression coefficients are also don't have the same interpretation. How do we deal with this?

- The idea is that we can create a new predictor from our already existing predictors which accounts for the interactions. To this end, consider the two predictor regression model written now as

$$Y = \beta'_0 + \beta'_1X_1 + \beta'_2X_2 + \epsilon.$$

I'm again using the ' to distinguish this model from an expanded model with a new predictor. The new predictor which accounts for the interaction will be $X_3 = X_1X_2$! Note that X_3 is going to be correlated with both X_1 and X_2 , but crucially it will not be perfectly correlated unless X_1 and X_2 are with each other. The new model is

$$Y = \beta_0 + \beta_1X_1 + \beta_2X_2 + \beta_{1,2}X_3 + \epsilon$$

Or written as a conditional distribution

$$Y|X \sim \text{Normal}(\beta_0 + \beta_1X_1 + \beta_2X_2 + \beta_{1,2}X_3, \sigma^2).$$

- Now how do we interpret the regression coefficients in the new model? Well, let's start by writing out the regression coefficient for X_1 in the model with no interaction:

$$\begin{aligned}\beta'_1 &= \mathbb{E}[Y|X_1 = x + 1, X_2] - \mathbb{E}[Y|X_1 = x, X_2] \\ &= (\beta_1(x + 1) + \beta_2X_2 + \beta_3(x + 1)X_2) - (\beta_1x + \beta_2X_2 + \beta_{1,2}xX_2) \\ &= \beta_1 + \beta_{1,2}X_2\end{aligned}$$

Here I've replaced X_3 with X_1X_2 . In words, $\beta_{1,2}$ is the additional difference in the regression slope between Y and X_1 with all other predictors fixed when X_2 is changed by one unit. If $\beta_{1,2}$ is positive (resp. negative) then increasing X_2 has a tendency to increase (resp. decrease) the difference between the conditional averages of Y , thus, the additional predictor $X_3 = X_1X_2$ allows us to capture an interaction wherein the association between X_1 and Y depends on X_2 . We can do the same calculation with X_1 fixed and X_2 changed to obtain

$$\mathbb{E}[Y|X_1 = x + 1, X_2] - \mathbb{E}[Y|X_1 = x, X_2] = \beta_2 + \beta_3X_1$$

Now as always, let's make sure we understand how this translates into code and also I think the visualizations in this example are helpful.

Example 1 (Visualizing interactions). Consider the regression model with interactions term $X_3 = X_1X_2$ defined above, assuming that

$$X_1 \sim \text{Normal}(0, 1)$$

$$X_2 \sim \text{Bernoulli}(1/2)$$

Assume $\beta_1 = 1.2, \beta_2 = 0$ and $\beta_3 = -1$.

Question: What are the slopes of $E[Y|X_1, X_2 = 0]$ vs. X_1 and $E[Y|X_1, X_2 = 1]$ vs. X_1 ? Plot these in colab.

Example 2 (Interaction in test score model). Here we once again consider the test score data with high school education and IQ as predictors. In particular, we will fit the model

$$Y = \beta_0 + \beta_{hs}X_{hs} + \beta_{iq}X_{iq} + \beta_{iq,hs}X_{iq}X_{hs}$$

Question:

- (a) What are the values and interpretations of the regression coefficients in the new model?
- (b) What do the results tell us about the “effect” of IQ on test scores and how it is related to high school education?

Solutions: The output of statsmodels is

```
> =====
>               coef      std err          t      P>|t|      [0.025      0.975]
> -----
> const         -11.4820      13.758       -0.835      0.404     -38.523      15.559
> mom_hs         51.2682      15.338        3.343      0.001      21.122      81.414
> mom_iq          0.9689       0.148        6.531      0.000        0.677        1.260
> hs_iq          -0.4843       0.162       -2.985      0.003      -0.803      -0.165
```

(a) The values and interpretations are as follows

- $\beta_0 \approx -11$ does not have a clear interpretation since it doesn't make sense to have 0 IQ.
 - $\beta_{iq} \approx 1$ is the expected signed difference in scores between students whose IQ differs by 1 point and whose mothers did not go to high school.
 - $\beta_{hs} \approx 51$ is the expected difference between scores of students whose mothers did and did not go to high school when IQ is zero – **much like β_0 , this doesn't really make sense, so we don't have a clear interpretation of β_{hs}**
 - $\beta_{hs,iq} \approx -0.5$ is the expected increase between the group of students whose mothers attended high school and those that did not in the difference between test scores of students whose mothers iq differs by one point. In other words, it is the increase in regression slope of scores vs. iq between the two high school groups. It is easier to interpret $\beta_{iq} + \beta_{hs,iq}$ which is the expected signed difference in scores between students whose IQ differs by 1 point and whose mothers did attend high school (this is the same as the interpretation of β_{iq} for student's whose mothers attended high school).
 - It appears that the association between IQ and test scores is weaker among students whose mother's attended high school. This makes intuitive sense. If someone attended high school the subtle differences in cognitive ability that IQ is supposed to measure are not as relevant.
- In the example above we saw that adding an interaction term can make the interpretation of the regression coefficients a bit clunky, even those regression coefficients that are not involved in an interaction. In particular, we ran into the problem that the interpretation of β_{iq} was lost. To remedy this, we can define the centered predictor

$$\Delta_{iq} = X_{iq} - \bar{X}_{iq}$$

and now we fit the model

$$Y = \beta_0 + \beta_{hs}X_{hs} + \beta_{iq}\Delta_{iq} + \beta_{hs-iq}X_{hs}\Delta_{iq} + \epsilon$$

Now β_1 has the interpretation as the average difference in scores among students whose mothers have the average iq. You can also standardize the predictor to get around the awkward interpretation of the regression coefficients.

3.1 Interactions

- We address the question of how we might identify when it is appropriate to add an interaction term to a model. In this case of two predictors we can easily visualize the data by making plots of the Y vs. X_1 slopes for different X_2 values (especially when one predictor is binary). With many predictors, it becomes less clear what plot might reveal some hidden interaction terms. We could of course go on adding every possible interaction term, but with many predictors this becomes in practical and leads to overfitting (as we will soon discuss).

- The basic idea of residual plots is that by plotting the difference between the observed y values and the prediction of the $E[Y|X]$, or

$$r_j = Y_j - \sum_i^K \hat{\beta}_i X_{i,j}$$

where $X_{i,j}$ is the value of predictor i at the j th data point. Our goal is to plot r_j in a such a way that disagreement between our model and the data is revealed by this plot. In the instance of a single-predictor, we can simply plot r_j as a function of the predictor X . If we notice that the residuals do not appear to follow a normal distribution, or that the variance and mean change, then we should be skeptical. The question is: **When we have multiple predictors, what do we plot on the horizontal axis?** The answer is to plot r_j as a function of the predictors value of $E[Y|X]$; that is r_j vs. $\hat{y}_j = \sum_i^K \hat{\beta}_i X_{i,j}$. The following examples is supposed to help us understand why.

Example 3 (Residual plots). *Consider the regression model with two predictors and suppose the true parameter values are $\beta_0 = 0, \beta_1 = 0.2, \beta_2 = 20$ and $\sigma = 1$.*

Question: Compare two ways of plotting the residuals

- Plot r_j as a function of $\sum_i^K \hat{\beta}_i X_{i,j}$.*
- Plot r_j as a function of Y_j . Why does there appear to be a bias towards high values of r_j for large Y_j that is not present in the first plot?*

Solution: See colab notebook

- Let's take a closer look at the mathematics underlying the residual plot. Note that

$$(1) \quad r_j \approx Y_j - \underbrace{E[Y|(X_1, \dots, X_K) = (X_{1,j}, \dots, X_{K,j})]}_{\approx \hat{y}_j}$$

where K is the number of predictors. Thus, the distribution of r_j is approximately

$$r_j \sim \text{Normal}(0, \sigma^2).$$

This tells us how the points should be distributed in the vertical direction. It **does not** say anything about the distribution of points in the horizontal direction, which is determined by the distribution of the predictors. Therefore, we expect a plot which is symmetric around the line r_j for all values of \hat{y}_j (our predicted values of Y), but any distribution in the horizontal direction is okay.

Now compare this to what would happen if we plotted Y_j on the horizontal axis, not \hat{y}_j . In this case, based on Equation 1 r_j and Y_j are correlated. This is why we see a bias of the residuals for small/large Y_j in Example 3.

4 Non-linear models

- Here, we discuss how to build more complex models and directly access their predictive power on out-of-sample data. In the context of interactions, we already saw how a model can be extended by defining a new predictor $X_3 = X_1 X_2$. The more general idea that we can define a new predictor which is a function of the other predictors allows us to develop very complex and flexible models which nonetheless can be analyzed within linear regression framework. Here, we will formalize this, beginning with the case of a single predictor.

In general, the linear regression framework allows us to fit models of the form consider the model

$$(2) \quad Y|X \sim \text{Normal}(f(X), \sigma^2)$$

provided we can express $f(X)$ as a linear combinations of nonlinear functions of X . What I mean by this is that we can find function $\phi_1(x), \dots, \phi_K(x)$ such that

$$f(X) = \sum_{i=1}^K \beta_i \phi_i(X)$$

The functions $\phi_i(X)$ are often referred to as basis functions, or features in machine learning lingo. We can think of each $\phi_i(X)$ as a new predictor.

Example 4 (Simulating a nonlinear model). *Consider the conditional Gaussian model given in Equation 2 with*

$$f(x) = \beta_0 + \beta_1 x + \beta_2 x^2$$

This is simply a linear model once we define the new predictors

$$X_1 = \phi_1(X) = X$$

$$X_2 = \phi_2(X) = X^2.$$

Question:

- Generate data from this model
- Fit the data to the model using `statsmodels`.

Solution: See colab notebook




Figure 1. An illustration of how a nonlinear dependence on our predictor can be incorporated into the linear modeling framework by adding a feature.

- Which functions ϕ should we use? The answer of course depends on the problem at hand. For example, we might know something about the physics of the data we are modeling. In some cases, we may select the ϕ so that the parameters a_i have clear interpretations (as is the case in linear regression model). The following illustrates such an example.

Example 5 (Mauna Kea Data). Solution: See colab notebook

- We now understand that within the linear regression framework we can fit very complicated nonlinear data. However, we need to be careful when adding new features to the model. The following example will illustrate how it is possible to “overshoot” and make our too complex. In the next set of notes we will dive much deeper into this, bridging the gap between machine learning and statistics.

Example 6 (Models with very large numbers of parameters). *Consider the model*

$$Y|X \sim \text{Normal}(f(X), \sigma^2)$$

$$f(X) = 1 - X^2 + 0.9X^3$$

This is a linear regression model with the features $X_1 = X^2$ and $X_2 = X^3$. We will pretend we have some data from this model, but not only do we lack knowledge of the coefficients, we in-fact lack knowledge of what the function $f(X)$ is. This, we are going to fit the data to a polynomial regression model

$$f(X) = \beta_0 + \beta_1 X + \beta_2 X^2 + \dots + \beta_K X^K$$

The question is which value of K to select, since we don't know a-priori that $K = 3$.

Question: *How does the predicted function $f(x)$, which we will denote $\hat{f}(x)$ compare to the true function? How does this depend on the number of features we include in our model?*

5 Orthogonality and Fourier analysis

- It is often to our advantage to select basis function ϕ_i so that our predictors are not correlated. That is, if $\phi_i(X)$ and $\phi_j(X)$ are thought of as random variables (the randomness comes from X) then we want them to be uncorrelated.
- Throughout our discussion, we will assume that $E[\phi_i(X)] = 0$. We can easily make this so by subtracting the means of our features. Therefore, we would like to find ϕ_i such that

$$(3) \quad E[\phi_i(X)\phi_j(X)] = 0 \quad i \neq j.$$

We this holds for some distribution of X , we say that ϕ_i and ϕ_j are orthogonal with respect to this distribution.

Example 7 (Sin series). *Let us now work with the particular example where*

$$(4) \quad X \sim \text{Uniform}(-L, L).$$

Even if our X points are not random, but are say evenly spaced on the interval $[-L, L]$, a random sample from the uniform distribution will be statistical similar to some randomly selected evenly spaced X points. Thus, we can think of the uniform distribution as approximating the spread of our X data with a random variable.

Now consider the basis functions

$$\phi_j(x) = \sin\left(\frac{\pi x j}{L}\right).$$

Question: *Show using simulations that ϕ_j are orthogonal with respect to the uniform distribution on $[-L, L]$ (Equation 4).*

Solution: *See colab notebook.*

- Notice that if ϕ_j are orthogonal, then we can express the regression coefficients as

$$\beta_j = \frac{\text{cov}(Y, \phi_j(X))}{\text{var}(\phi_j(X))} = \frac{E[Y\phi_j(X)]}{E[\phi_j(X)^2]}$$

and hence our fitted coefficients from N data points can be approximated as

$$(5) \quad \hat{\beta}_j \approx \frac{\sum_{i=1}^N \phi_j(X_i) Y_i}{\sum_{i=1}^N \phi_j(X_i)^2}$$

This suggests $\hat{\beta}_j$ should depend only weakly how many of the features ϕ we have included in our model! This is consequence of orthogonality. In contrast, in the earlier example where we used polynomial features adding a new predictor, say X^4 , would dramatically change the fitted value of β_2 . However,

- Now let's return to Example 7. A limitation of this model is that it only permits us to model "odd functions" that is, functions for which $f(x) = -f(x)$, as each ϕ_j has this property. It is therefore desirable to add additional features which permit this, but we'd like them to still be orthogonal. To this end, we introduce a new set of features to the model given by the cosine functions:

$$\cos\left(\frac{\pi x j}{L}\right).$$

Our new model, called a Fourier series, is given by

$$f(x) = \beta_0 + \sum_{i=1}^K \beta_i \sin\left(\frac{\pi i x}{L}\right) + \alpha_i \cos\left(\frac{\pi i x}{L}\right)$$

where we are using α_i to represent the coefficients of the cosine terms. Fourier series are one of the most important models in data science and engineering. It can be proved that as $K \rightarrow \infty$ we can approximate essentially ANY function with a series of this form.

Example 8 (Working with Fourier series in python). *Suppose we have data from the model with*

$$f(x) = \sin(3\pi x) + \sin(10\pi x)$$

where $X_i = (i - 1)/N$ (meaning the X points are evenly spaced in $[0, 1]$).

Question: For different values of σ^2 , generate data from this model and fit it to a Fourier series with $K = 100$ terms.

Solution: See colab notebook.

- The process of computing the coefficients $\hat{\beta}_j$ and $\hat{\alpha}_j$ for the Fourier series is called a Discrete Fourier Transform. Usually DFT refers to the cases where the X_i are equally spaced. In this case, the orthogonality condition (Equation 3) is true in "data world", not just "math world" – what I mean by this is that for equally spaced data points.

To be precise, consider the predictor data on the interval $[0, L]$ given by

$$X_i = \frac{L(i - 1)}{N - 1}$$

for $i = 1, \dots, N$. Thus $X_1 = 0$, $X_2 = L/N$, $X_3 = 2L/N$ and $X_N = L$. We could generate these points with `np.linspace`. The following theorem tells us that the empirical, or sample covariance between the sin features, is exactly zero for these predictors.

Theorem 1.

$$\sum_{i=1}^N \sin\left(\frac{2\pi j X_i}{L}\right) \sin\left(\frac{2\pi k X_i}{L}\right) = 0 \quad k \neq j$$

- Often we want to summarize how different frequencies are represented in our data, but we don't particularly care about whether they come from the sin or cos terms. To achieve this, one uses the power spectrum density, also known as the periodogram,

$$P_j = \beta_j^2 + \alpha_j^2.$$

The power spectrum density is a fundamental object in signal processing, and it essentially tells us how "wobbly" a signal is.

Example 9 (Periodogram). Question: Compute the periodogram of the data generated in Example ?? and confirm the same periodogram can be generated with the `periodogram` function from the `scipy.signal` library. This is neat and not often made point, that this fundamental structure from signal processing is in fact coming from fitting a linear regression model with least square!!

Solution: See *colab notebook*.