EVALUATING MODEL PREDICTIONS, FOURIER SERIES

Contents

1 Learning objectives

- Cross validation and how to use it to perform model selection.
- Bias-variance decomposition of mean-squared error, and why it explains the U-shaped test error curve.
- Fourier series as linear regression models. Orthogonality and why it is important.

2 Cross validation

- We saw from the previous example that R^2 is not very helpful if we want to decide how complex we'd like to make our model, since it is possible to explain all the variation in the Y values with a model which is too complex. In order to address this, we take the approach of cross validation. The basic idea of cross validation is to break our data up into two subsets: training set, which we use to fit the model, and a testing set, which we compare to the predictions of the fitted model.
- **Some Notation**
  - I will use D to refer to our entire data set:
  $$D = (Y, X) = \{(X_1, Y_1), \ldots, (X_N, Y_N)\}$$
  - $D^{\text{train}} = (Y^{\text{train}}, X^{\text{train}})$ and $D^{\text{test}} = (Y^{\text{test}}, X^{\text{test}})$ will represent the subsets of the data to be used to training (that is, fitting) the model and testing the fit respectively. We will assume that
  $$D = D^{\text{train}} \cup D^{\text{test}}$$
  Let N_{train} and N_{test} denote the number of points in each group.
  - Let $\hat{y}(x, D)$ the prediction of $E[Y|X = x]$ using the fitted coefficients based on a data set D; that is
  $$\hat{y}(x, D) = \sum_{i=1}^{K} \hat{\beta}_i \phi_i(x)$$
  where $\hat{\beta}$ are the (usually least squares) fitted coefficients using the data in D.
- Now let us define the training error
$$\epsilon^2_{\text{train}} = \frac{1}{N_{\text{train}}} \sum_{i=1}^{N_{\text{train}}} (\hat{y}(X_i^{\text{train}}, D^{\text{train}}) - Y_i^{\text{train}})^2$$
where the average is taken over different replicates of our data and $\hat{Y}_i$ is our prediction of $E[Y|X]$. Not that the training error tells us about how well our model does at predicting the same data we used to fit it, and therefore not surpisingly, it is closely related to R^2:
$$R^2 \approx 1 - \frac{\epsilon^2_{\text{train}}}{\text{var}(Y_{\text{train},i})}.$$

In order to see how well our model does at predicting the points we did NOT use to fit it, we introduce the test error:

$$\epsilon^2_{\text{test}} = \frac{1}{N_{\text{test}}} \sum_{i=1}^{N_{\text{test}}} (\hat{y}(X_i^{\text{test}}, D^{\text{train}}) - Y_i^{\text{test}})^2$$

Note that the fitted coefficients used to compute $\hat{y}_i(X^{\text{test}})$ come from fitting the model to the training data, even though we are evaluating $\hat{y}$ at the test points.

Example 1 (Cross validation on polynomial model)**.** *Consider example 6 from the previous weeks notes.*

Question: Plot the training and test error as a function of the number of parameters.

Solution: See colab notebook.

3 Bias, variance and overfitting

- In order to understand U-shaped curve seen in Example 1, we first need to introduce some more general terminology and notation for talking about estimators (of which $\hat{y}$ is an example). To this end, we introduce the mean-squared error of an estimator $\hat{\theta}$ of some quantity θ. θ could be a parameter, or it could be a value of a function, such as $f(x)$, that we would like to predict.

$$\text{MSE}_{\hat{\theta}} = E\left[(\hat{\theta} - \theta)^2\right] \tag{1}$$

  For now, let's just think of $\hat{\theta}$ as any estimator.
- The following theorem is the key result which will allow us to understand the U-shaped curve.

  Theorem 1 (Bias variance decomposition)**.**

$$\text{MSE}_{\hat{\theta}} = \text{var}(\hat{\theta}) + E\left[\hat{\theta} - \theta\right]^2 \tag{2}$$

  Proof. Using the definition of variance

$$\text{var}\left(\hat{\theta} - \theta\right) = E\left[(\hat{\theta} - \theta)^2\right] - E\left[\hat{\theta} - \theta\right]^2.$$

  Since θ is a constant, $\text{var}\left(\hat{\theta} - \theta\right) = \text{var}\left(\hat{\theta}\right)$, so rearranging terms yields the result. □

  Now let's break this result down a bit.
  - The first term in Equation 2 is simply the variance, which in the case of an estimator of a parameter we can recognize as the squared standard error. This is tell us how much variation there is in our estimate from replicate of replicate. You should recognize that a consistent estimator is exactly one for which this quantity vanishes when N is very large.
  - The second term is what we will define as the squared of the bias:

$$\text{Bias}_{\hat{\theta}} = E[\hat{\theta} - \theta].$$

    This tell us whether our estimate will on average give us the correct value of θ. You should recognize that an unbiased estimator is exactly one for which this quantity is zero!
- Now let's return to working in the context of a model of the form

$$Y|X \sim \text{Normal}(f(X), \sigma^2)$$

  where $f(X)$ can be written as a linear combination of features

$$f(X) = \sum_{i=1}^{N} \beta_i \phi_i(X).$$

  Letting $\hat{\theta} = \hat{y}(x, D)$, applying Equation 1 yields

$$\text{MSE}_{\hat{y}(x,D)} = E\left[(\hat{y}(x, D) - f(x))^2\right]$$

  where the average is taken over replicates of our data D. This is a measurement of our ability to predict $f(x)$, which is exactly what the test error seeks to measure.

- To more closely relate the bias-variance tradeoff to the test error, we will use a slightly different definition of MSE, which is

$$\widetilde{\text{MSE}}_{\hat{y}(x,D)} = E\left[(\hat{y}(x,D) - Y(x))^2\right] \tag{3}$$

where Y is a sample from $Y|(X = x)$. You can show (see Exercise) that

$$\widetilde{\text{MSE}}_{\hat{y}(x,D)} = \sigma^2 + \text{var}(\hat{y}(x,D)) + E\left[\hat{y}(x,D) - f(x)\right]^2$$

The new term, σ^2, captures the fact that, unlike the deterministic term f, we can NEVER predict the *random* variable Y exactly. You should be able to justify that

$$\epsilon^2_{\text{test}} \approx E\left[\widetilde{\text{MSE}}_{\hat{y}(X,D^{\text{train}})}\right]$$

- The important thing to recognize is that a complicated model will tend to have a higher variance, since it will be able to change more in response to data. Meanwhile, a very simple model will tend to have a higher bias.

Example 2 (Plotting Bias and variance)**.** *Let's continue with Example 1 and try to illustrate the bias variance tradeoff by computing these (or really approximations to them) separately.*

Solution: See colab notebook.

- We now summarize some observations we've made thus far:
  - In a regression model, as we add more coefficients, eventually the model becomes more and more "flexible", in the sense that is can describe more different types of data sets. Here, by "describe", we mean that it can be fit to those data sets with a small value of R^2 or training error. With enough features, a model can perfectly interpolate between the data points, meaning $\hat{y}(X_i, D) = Y_i$ for each data point (X_i, Y_i) when the model is fit on all our data points.
  - Unlike the training error and R^2, the test error, ϵ^2_{test} does not simply increase as we make the model more complex, rather it has a U-shape. Thus, there is a optimal model size at which our model's predictions of new data points – that is, data points outside the set we used to fit it – is best.
  - The U-shape can be understood in terms of bias and variance. We know this because the mean-squared error, which is the "math world" version of ϵ^2_{test}, can be decomposed into a bias and variance term. The variance term tells us how variable the predictions of our model will be when we fit it to different training sets, while the bias tells us how much its predictions. will differ, on average, from real data.

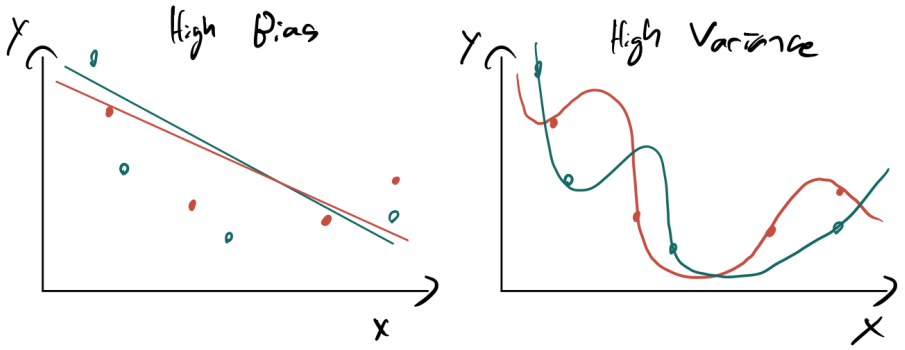

Figure 1. Bias and variance illustrated with fits to two different datasets (blue and red) drawn the from the same distribution